

Word embeddings

Based on lecture by Murat Apishev



- **Word embeddings idea.
One-hot vectors. SVD**
- **Word2vec model, training methods**
- **Word2vec training optimization: hierarchical softmax and SGNS**
- **GloVe model**
- **FastText model, Hashing Trick**
- **Word embeddings models' quality evaluation**



**Word embeddings idea.
One-hot vectors. SVD**



Input data may have different format



Visual content

- Images
- Video



Texts

- Unstructured documents
- HTML / XML



Structured documents

- Tables with features




Signals

- Audio: music / speech
- Other signals




Input data may have different format



Result



Input data may have different format



Word embeddings

"I love watching TV series"

text

"I" "watching" "series"

"love" "TV"

word

$\begin{bmatrix} 123 \\ 456 \\ 12 \\ \dots \\ 89 \end{bmatrix}$	$\begin{bmatrix} 23 \\ 372 \\ 8 \\ \dots \\ 83 \end{bmatrix}$	$\begin{bmatrix} 16 \\ 124 \\ 76 \\ \dots \\ 29 \end{bmatrix}$	$\begin{bmatrix} 2 \\ 12 \\ 299 \\ \dots \\ 65 \end{bmatrix}$	$\begin{bmatrix} 177 \\ 6 \\ 504 \\ \dots \\ 304 \end{bmatrix}$
---	---	--	---	---

embeddings



What kind of embeddings do we want?

Different words  Different embeddings

Ecology
Love

Close words  Close vectors

Love
Adore



What does «close» mean?

Semantic closeness

«Usual» word closeness

Examples:

- computer
- laptop
- PC

Closeness of embeddings

$\text{sim}(w, v)$

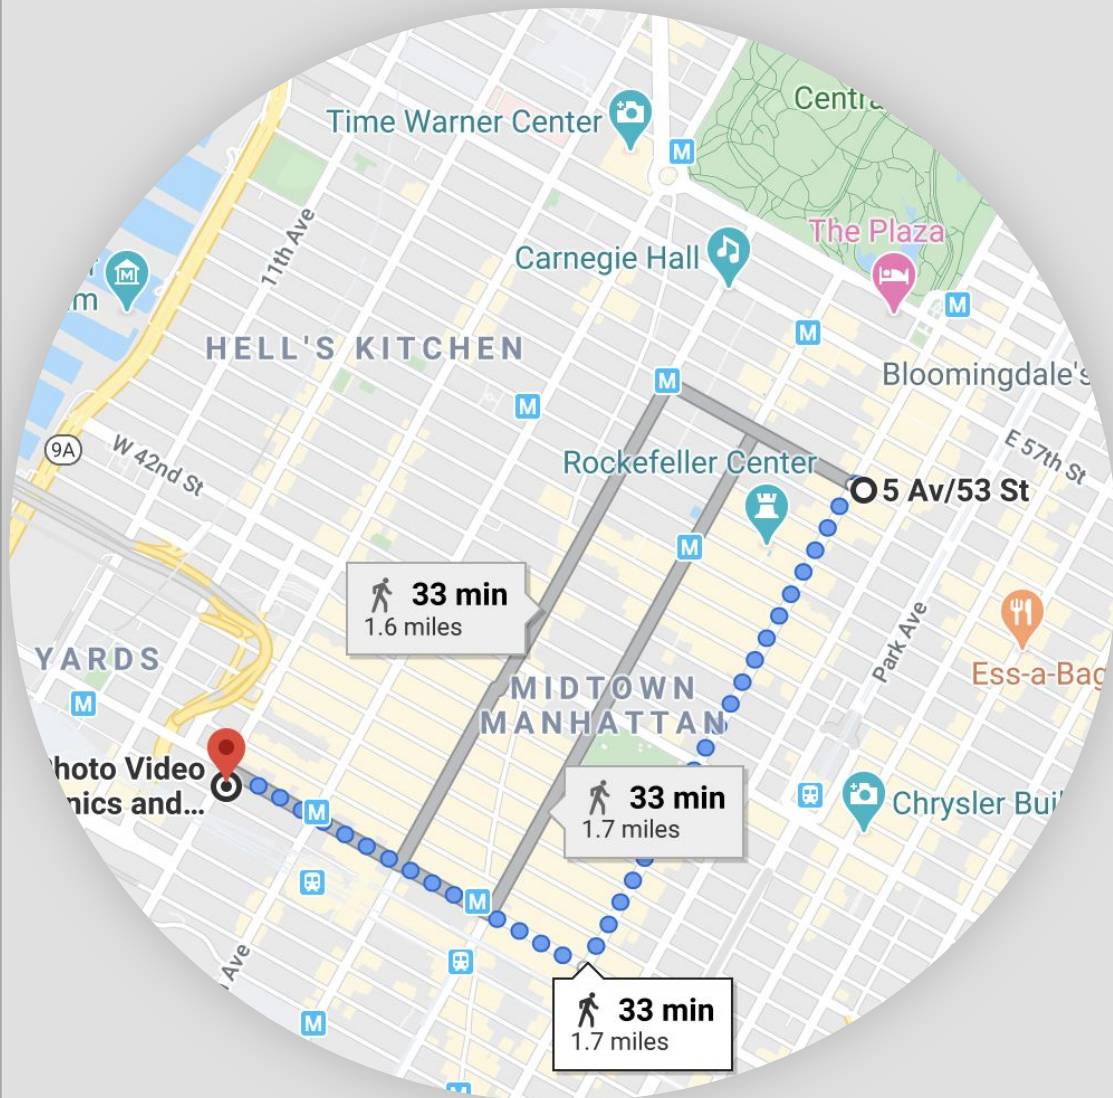
where $\text{sim}()$ can be:

- Manhattan distance
- Euclidean distance
- Cosine distance



Manhattan distance

$$d(p, q) = \sum_{i=1}^n |p_i - q_i|$$

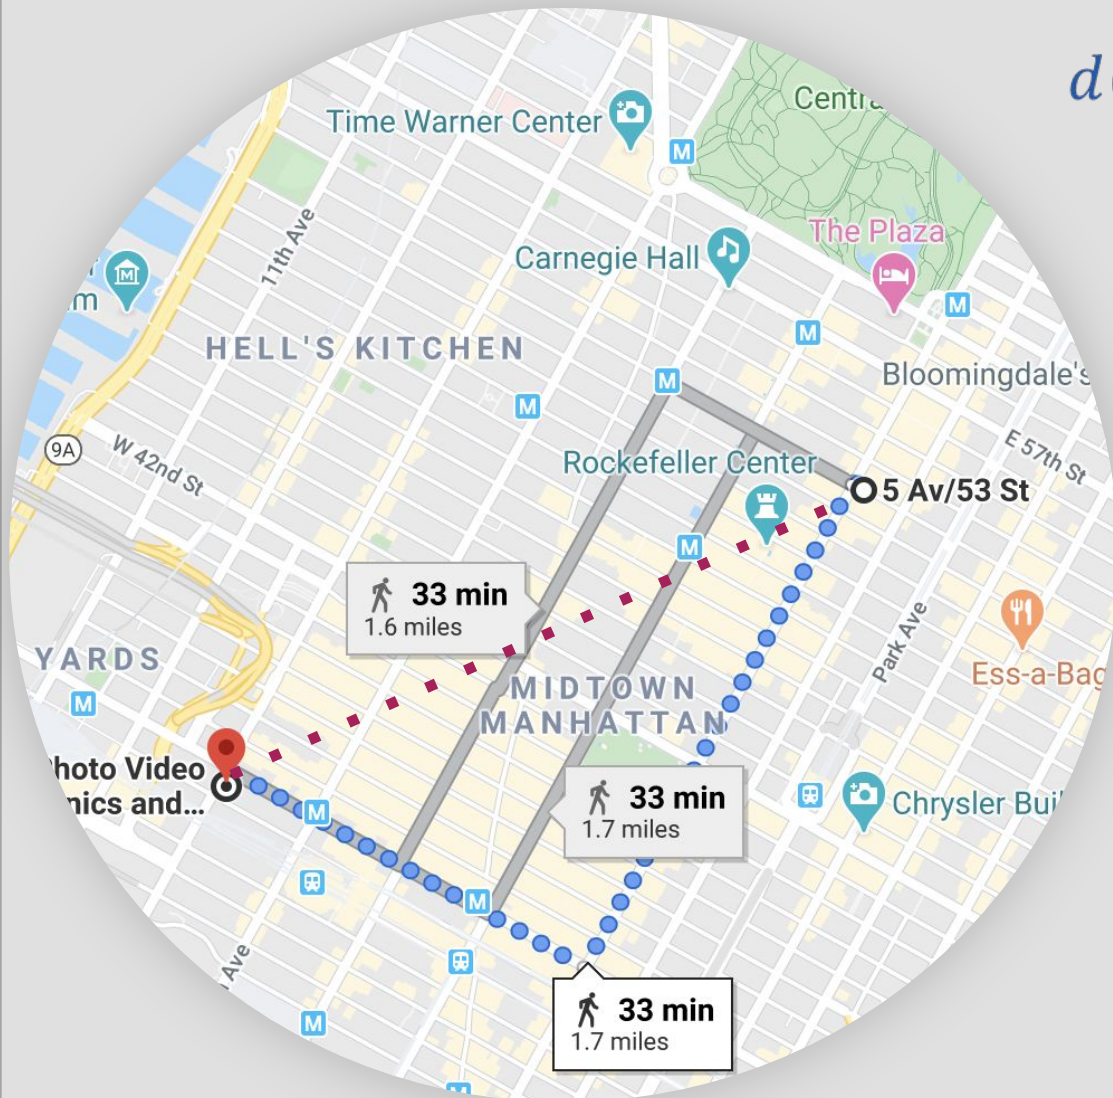


google.com/maps



Euclidean distance


$$d(p, q) = \sqrt{\sum_{i=1}^n |p_i - q_i|^2}$$




Cosine distance

$$p \cdot q = \|p\| \|q\| \cos(\theta)$$

- $$\cos(\theta) = \frac{p \cdot q}{\|p\| \|q\|} = \frac{(\sum_{i=1}^n p_i q_i)}{\sqrt{\sum_{i=1}^n p_i^2} \sqrt{\sum_{i=1}^n q_i^2}}$$



$$\theta \approx 45^\circ$$
$$\cos(\theta) \approx 0.7$$

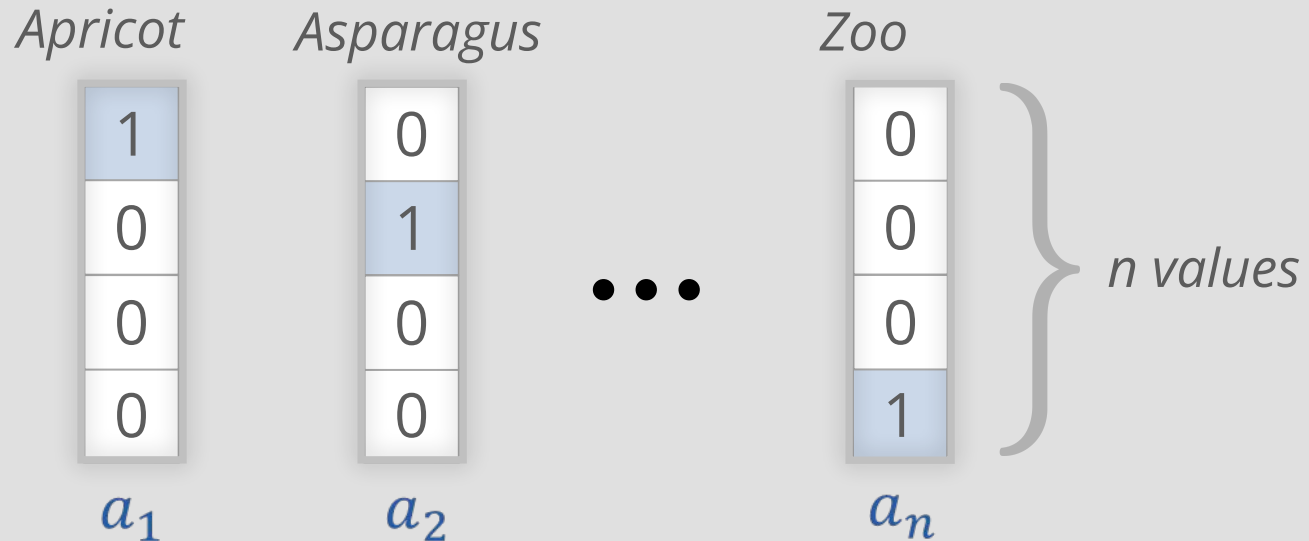


$$\theta \approx 7^\circ$$
$$\cos(\theta) \approx 0.99$$



One-hot encoding

Suppose we have vocabulary V , $|V| = n$



Advantages

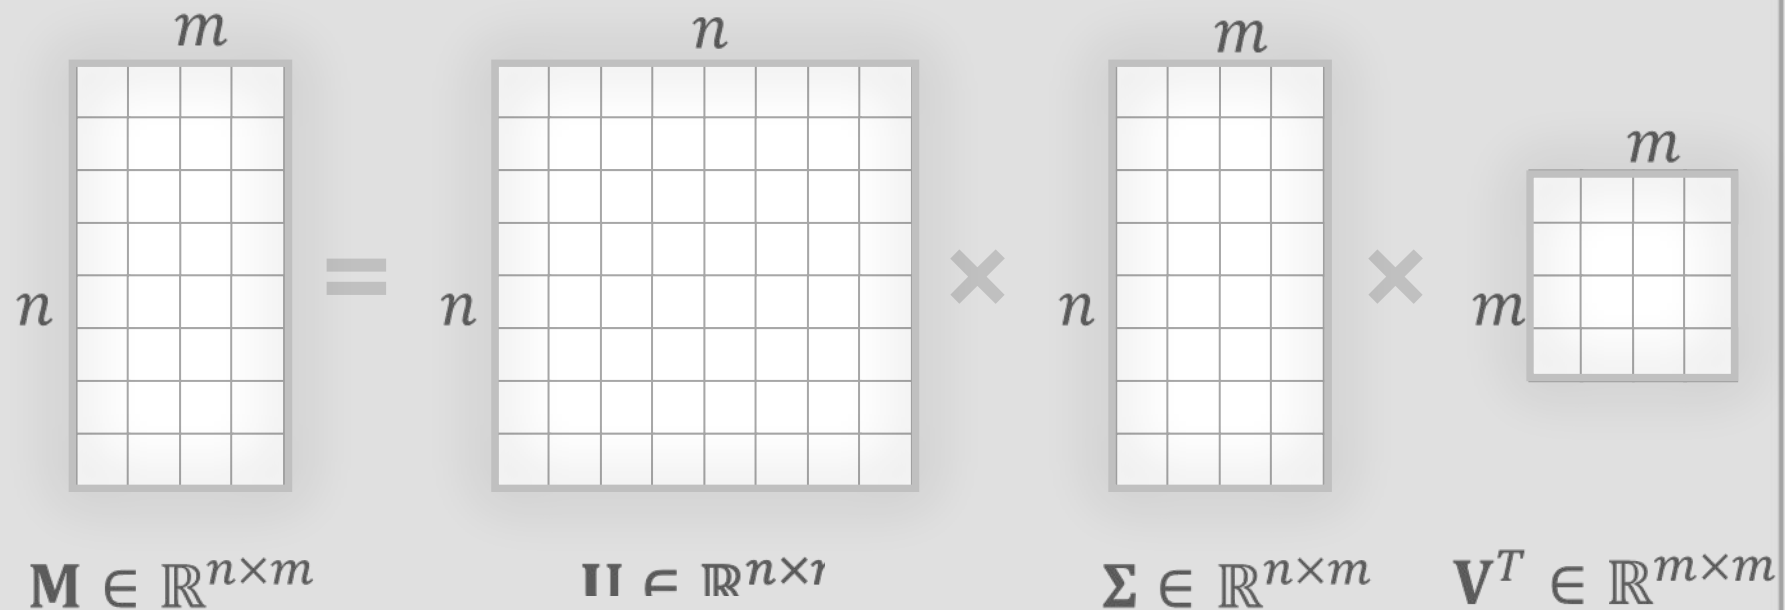
- Simple way to obtain embeddings for a set of words

Disadvantages

- The documents will have huge **unfixed length**
- Embeddings are **mutually orthogonal**



Singular Value Decomposition (SVD)


$$\mathbf{M} \in \mathbb{R}^{n \times m} = \mathbf{U} \in \mathbb{R}^{n \times n} \times \mathbf{\Sigma} \in \mathbb{R}^{n \times m} \times \mathbf{V}^T \in \mathbb{R}^{m \times m}$$



Singular Value Decomposition (SVD)

Suppose we have rectangular matrix $\mathbf{M} \in \mathbb{R}^{n \times m}$

The diagram illustrates the Singular Value Decomposition (SVD) of a matrix \mathbf{M} . It shows the equation $\mathbf{M} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$ using grid representations for each matrix. The first matrix, \mathbf{M} , is a blue grid with dimensions n (rows) and m (columns). The second matrix, \mathbf{U} , is a white grid with dimensions n (rows) and n (columns). The third matrix, $\mathbf{\Sigma}$, is a white grid with dimensions n (rows) and m (columns). The fourth matrix, \mathbf{V}^T , is a white grid with dimensions m (rows) and m (columns). The matrices are separated by an equals sign and multiplication symbols (\times). Below each grid is its corresponding mathematical notation and dimensions: $\mathbf{M} \in \mathbb{R}^{n \times m}$, $\mathbf{U} \in \mathbb{R}^{n \times n}$, $\mathbf{\Sigma} \in \mathbb{R}^{n \times m}$, and $\mathbf{V}^T \in \mathbb{R}^{m \times m}$.


$$\mathbf{M} \in \mathbb{R}^{n \times m} = \mathbf{U} \in \mathbb{R}^{n \times n} \times \mathbf{\Sigma} \in \mathbb{R}^{n \times m} \times \mathbf{V}^T \in \mathbb{R}^{m \times m}$$



Singular Value Decomposition (SVD)

Suppose we have rectangular matrix $\mathbf{M} \in \mathbb{R}^{n \times m}$

The matrix can be represented as $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$



Singular Value Decomposition (SVD)

- $\Sigma \in \mathbb{R}^{n \times m}$ is a diagonal matrix with non-negative values
- diagonal values of matrix Σ are singular values of matrix \mathbf{M}

The diagram illustrates the Singular Value Decomposition (SVD) of a matrix \mathbf{M} . It shows the equation:

$$\mathbf{M} \in \mathbb{R}^{n \times m} = \mathbf{U} \in \mathbb{R}^{n \times n} \times \mathbf{\Sigma} \in \mathbb{R}^{n \times m} \times \mathbf{V}^T \in \mathbb{R}^{m \times m}$$

Each matrix is represented by a grid of cells:

- $\mathbf{M} \in \mathbb{R}^{n \times m}$: An $n \times m$ grid with dimensions n (vertical) and m (horizontal) labeled above and to the left.
- $\mathbf{U} \in \mathbb{R}^{n \times n}$: An $n \times n$ square grid with dimension n labeled above and to the left.
- $\mathbf{\Sigma} \in \mathbb{R}^{n \times m}$: An $n \times m$ grid with dimensions n (vertical) and m (horizontal) labeled above and to the left. The first four rows and columns are highlighted in blue, showing a diagonal structure with 1s on the diagonal and 0s elsewhere.
- $\mathbf{V}^T \in \mathbb{R}^{m \times m}$: An $m \times m$ square grid with dimension m labeled above and to the left.

The matrices are connected by an equals sign and multiplication symbols (\times).



Singular Value Decomposition (SVD)

Columns **U** and **V** are left and right singular vectors of matrix **M**

The diagram illustrates the Singular Value Decomposition (SVD) of a matrix M . It shows the equation $M = U \Sigma V^T$ using grid representations for each matrix and their dimensions.

- Matrix M :** A grid with 8 rows and 4 columns, labeled with n on the left and m on the top. Below it is the notation $M \in \mathbb{R}^{n \times m}$.
- Matrix U :** A grid with 8 rows and 8 columns, labeled with n on the left and n on the top. Below it is the notation $U \in \mathbb{R}^{n \times n}$.
- Matrix Σ :** A grid with 8 rows and 4 columns, labeled with n on the left and m on the top. The diagonal elements are 1, 1, 1, and 1, with the rest being 0. Below it is the notation $\Sigma \in \mathbb{R}^{n \times m}$.
- Matrix V^T :** A grid with 4 rows and 4 columns, labeled with m on the left and m on the top. Below it is the notation $V^T \in \mathbb{R}^{m \times m}$.

The matrices are connected by an equals sign and multiplication symbols (\times) to show the decomposition: $M = U \Sigma V^T$.



Singular Value Decomposition (SVD)

Suppose we have a collection from m documents with n unique words

Then $\mathbf{M} \in \mathbb{R}^{n \times m}$ is a bag of words matrix

Let's apply SVD: $\mathbf{M} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$

The diagram illustrates the SVD decomposition of a matrix \mathbf{M} into three matrices: \mathbf{U} , $\mathbf{\Sigma}$, and \mathbf{V}^T .

Matrix \mathbf{M} is an $n \times m$ matrix, represented by a grid of 8 rows and 4 columns. Below it is the label $\mathbf{M} \in \mathbb{R}^{n \times m}$.

Matrix \mathbf{U} is an $n \times 1$ matrix, represented by a grid of 8 rows and 1 column. Below it is the label $\mathbf{U} \in \mathbb{R}^{n \times 1}$.

Matrix $\mathbf{\Sigma}$ is an $n \times m$ matrix, represented by a grid of 8 rows and 4 columns. The first 4 rows and 4 columns contain 1s, and the rest are 0s. Below it is the label $\mathbf{\Sigma} \in \mathbb{R}^{n \times m}$.

Matrix \mathbf{V}^T is an $m \times m$ matrix, represented by a grid of 4 rows and 4 columns. Below it is the label $\mathbf{V}^T \in \mathbb{R}^{m \times m}$.

The decomposition is shown as $\mathbf{M} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$, with multiplication symbols (\times) between the matrices.



Singular Value Decomposition (reduced SVD)

Pick top- k largest values of Σ

Ignore all other values and columns of matrix U - we obtain matrices U_k and Σ_k

The diagram illustrates the reduced SVD decomposition of a matrix $M \in \mathbb{R}^{n \times m}$. The matrix M is shown as a grid with n rows and m columns. It is decomposed into three matrices: $U_L \in \mathbb{R}^{n \times l}$, $\Sigma \in \mathbb{R}^{n \times m}$, and $V^T \in \mathbb{R}^{m \times m}$. The matrix U_L is a grid with n rows and l columns, where the first l columns are highlighted in blue. The matrix Σ is a grid with n rows and m columns, where the top l rows and the first l columns are highlighted in blue, forming a submatrix Σ_L . The matrix V^T is a grid with m rows and m columns, where the top l rows are highlighted in blue. The decomposition is shown as $M = U_L \Sigma V^T$.

$$M \in \mathbb{R}^{n \times m} = U_L \in \mathbb{R}^{n \times l} \times \Sigma \in \mathbb{R}^{n \times m} \times V^T \in \mathbb{R}^{m \times m}$$



Word embeddings with reduced SVD

We can use rows of matrix $\mathbf{U}_k \sqrt{\boldsymbol{\Sigma}_k}$ as word embeddings

The diagram shows the equation $\mathbf{E} = \mathbf{U}_k \sqrt{\boldsymbol{\Sigma}_k}$. Matrix \mathbf{E} is $n \times l$ with rows labeled 'Apricot', 'Asparagus', ..., 'Zoo'. Matrix \mathbf{U}_k is $n \times k$. Matrix $\sqrt{\boldsymbol{\Sigma}_k}$ is $k \times k$ and is shown as a 3x3 matrix with blue diagonal elements and zeros elsewhere.

$$\mathbf{E} \in \mathbb{R}^{n \times l} \quad \mathbf{U}_k \in \mathbb{R}^{n \times k} \quad \sqrt{\boldsymbol{\Sigma}_k} \in \mathbb{R}^{k \times k}$$

By the way:

для получения
векторов
раскладывать
можно не
только матрицу
«мешка слов»



Word embeddings with SVD

Improvements:

- Vectors have fixed size
- Vectors are no longer mutually orthogonal
- Semantic closeness is somehow taken into account

Nevertheless:

- Adding new words/documents requires new SVD calculation
- We need to operate with a huge BoW matrix
- Word embeddings are not that good



Main conclusions

- Word embeddings are used as features in NLP tasks
- Good word embeddings represent semantic closeness of words
- One-hot vectors can be useful but they are too sparse and mutually orthogonal
- SVD can produce word embeddings of fixed size that somehow represent semantics



Word2vec model



Distributive semantics hypothesis

- Words with similar meaning share similar context

«Today I ate **tasty**, **juicy** orange»
«This **apple** is so **sweet** and **juicy**»
«So **sweet** are the **apricots**, so **tasty**»

- Instead of frequency counters let's train a model to predict a word by its context (and vice versa)

- Harris Zellig. *Distributional structure // Word.* — 1954. — Vol. 10, no. 23. — Pp. 146–162.



Word2vec models: CBOW and skip-gram

- Continues Bag-of-Words

predict central word by its context



- Skip-gram

predict context by central word




- *Distributed Representations of Words and Phrases and their Compositionality.* / Tomas Mikolov, Ilya Sutskever, Kai Chen et al. // NeurIPS — 2013. — Pp. 3111–3119.



CBOW model

Suppose we have a collection with N unique words


To train a model we slide over text with a window of size $2C + 1$



step 1




CBOW model



step 2




CBOW model



step 3



CBOW model



step

4



CBOW model

I shape an apricot with friend

step 5



CBOW model

I share an apricot with friend

step


6



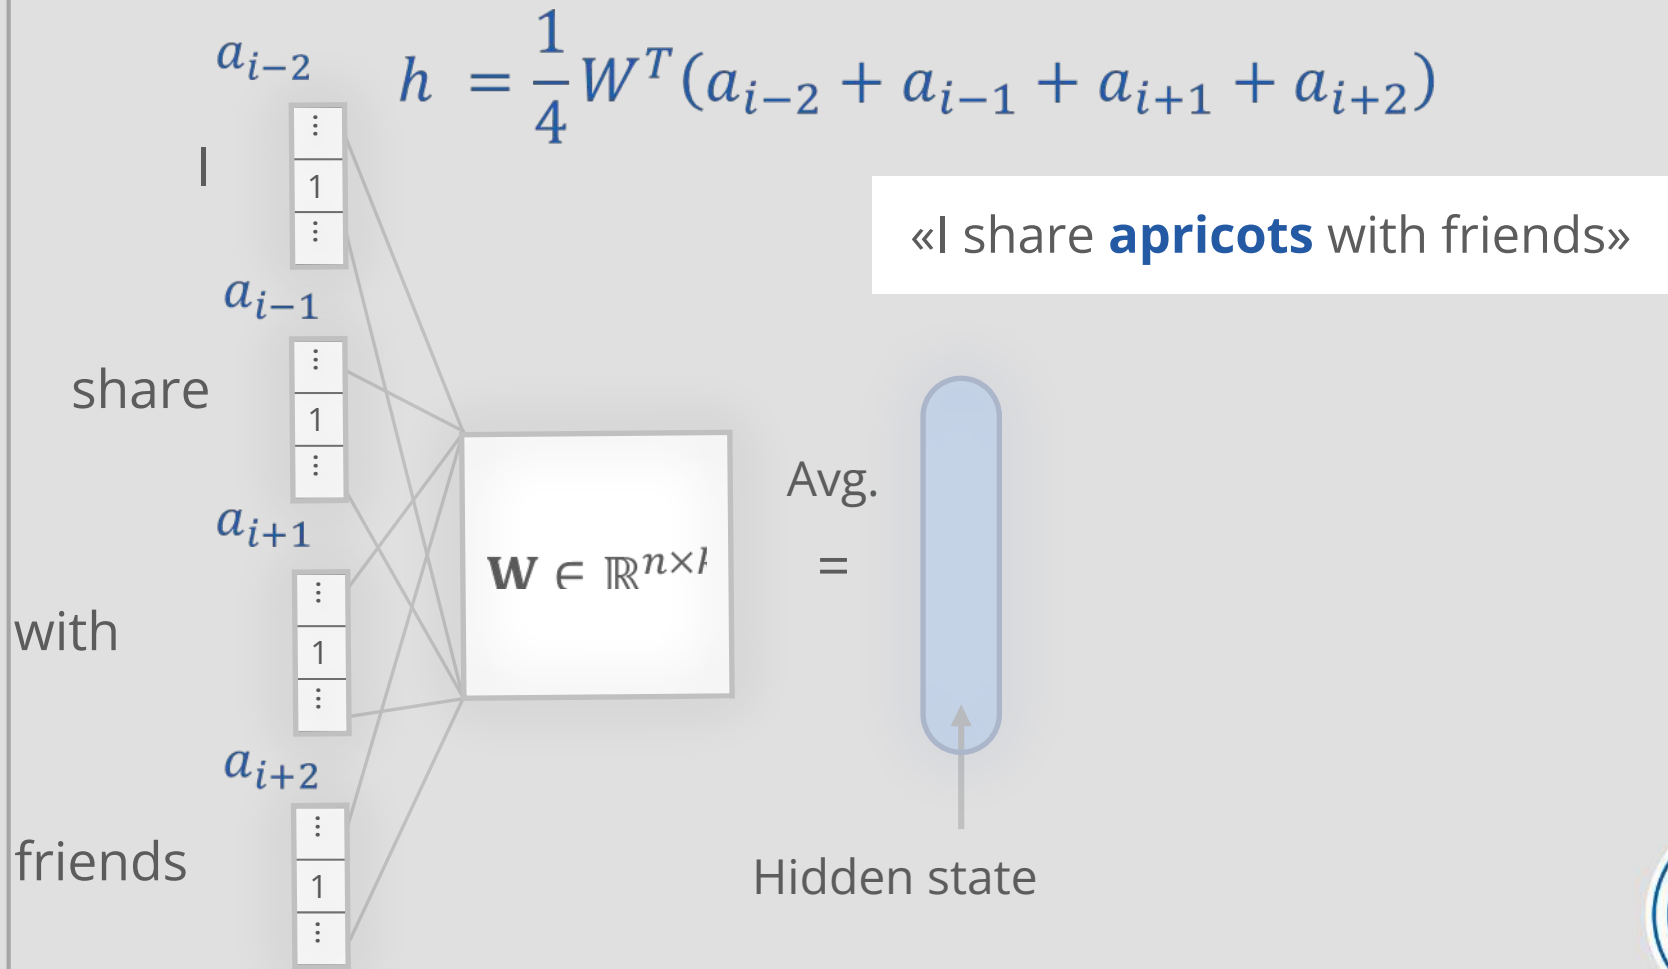
CBOW as neural network

Model - two-layer neural network

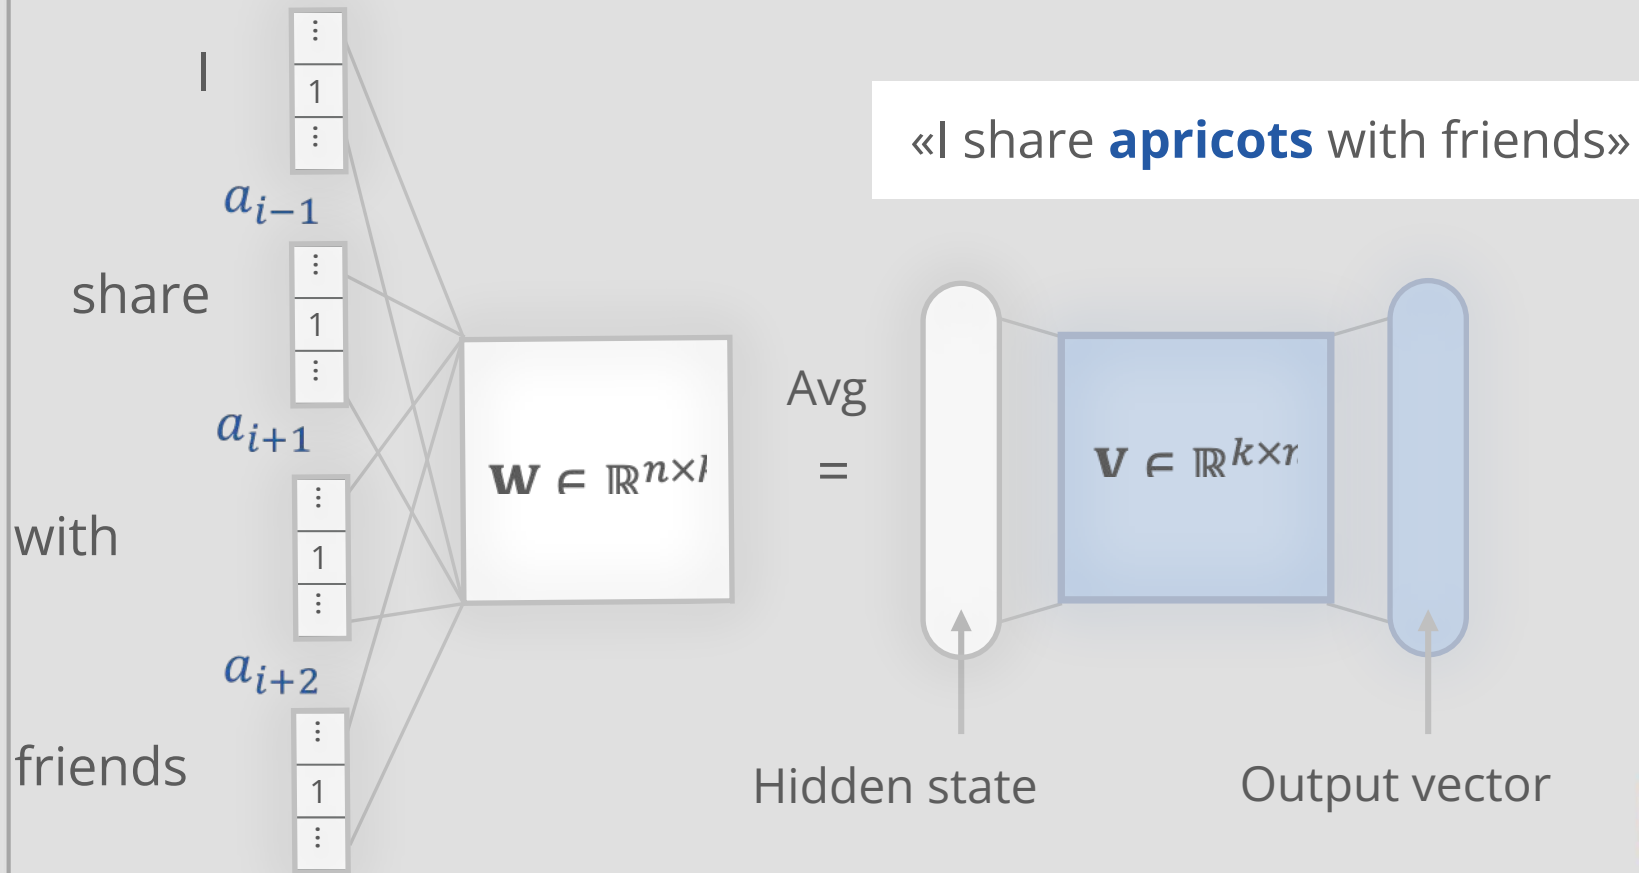
Input - $2C$ one-hot context vectors of size n



CBOW as neural network



CBOW as neural network



Loss function

- For window with index i predict word w_i by context c_i

$$\sum_{i=1}^N \log p(w_i | c_i) \rightarrow \max_{W, V}$$



Loss function

- For window with index i predict word w_i by context c_i

$$\sum_{i=1}^N \log p(w_i | c_i) \rightarrow \max_{w, v}$$

- Therefore model's output is a vector with n probabilities



Loss function

- For window with index i predict word by context

$$\sum_{i=1}^N \log p(w_i | c_i) \rightarrow \max_{W, V}$$

- Therefore model's output is a vector with n probabilities
- Therefore model's output is a vector with n probabilities

Softmax function:

$$\text{softmax}(b) = (z_1, \dots, z_n), \quad z_j = p(w_j | c_i) = \frac{e^{b_j}}{\sum_{k=1}^n e^{b_k}}$$



And the word embeddings?


- As a result of training we obtain two matrices: W and V
- Usually, rows of matrix W are used as word embeddings
- But both columns of V and the combination of two matrices can be used



Skip-gram as neural network

- Skip-gram model is arranged symmetrically

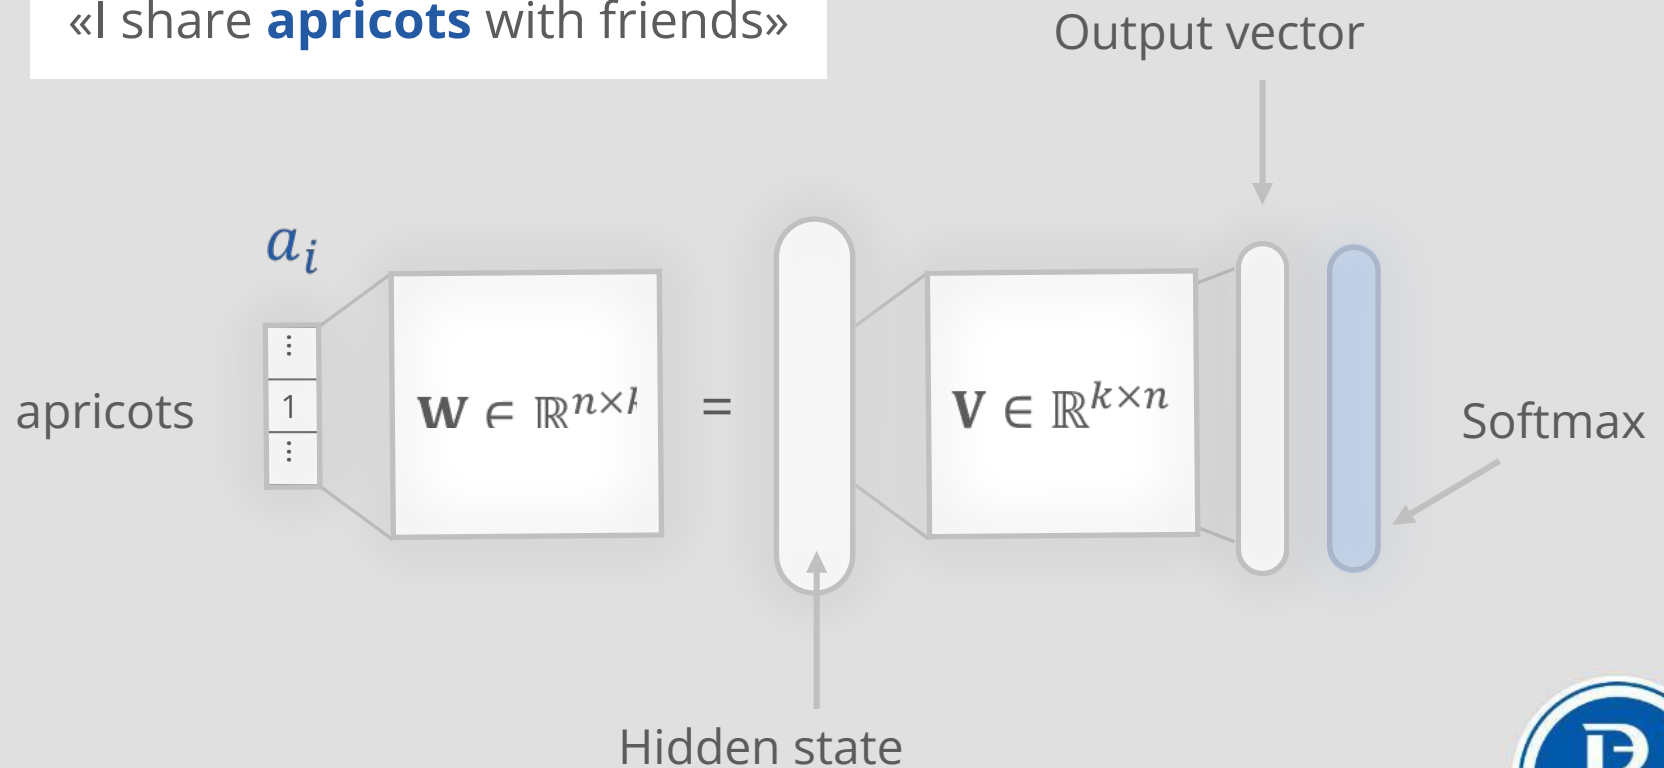
«I share **apricots** with friends»



Skip-gram as neural network

- Output — one probability distribution over words for the central word

«I share **apricots** with friends»



Loss function

- Output — one probability distribution over words for the central word
- $2C$ words from actual context should have maximal values
- Loss function:

$$\sum_{i=1}^N \sum_{j=-C, j \neq 0}^C \log p(w_{i+j}|w_i) \rightarrow \max_{W,V}$$



Main conclusions

- Word2vec models train word representations based on predictions, not on statistics
- There are two basic models: CBOW and Skip-gram
- In canonical implementation word2vec is a two-layer neural network, and its weights are the resulting word embeddings
- The quality of word2vec embeddings is better than SVD embeddings, and we don't need huge BoW matrices anymore



Word2vec training optimization: hierarchical softmax and SGNS



Word2vec neural approach disadvantages

- Training word2vec neural network is computationally difficult
- We need to calculate softmax ($O(n)$) and update a lot of parameters



Word2vec neural approach disadvantages

- Training word2vec neural network is computationally difficult
- We need to calculate softmax ($O(n)$) and update a lot of parameters
- In practice, word2vec is trained with optimization methods



Hierarchical softmax

- We still have a **fully connected neural network**
- The only thing that differs is **softmax** calculation

- *Mnih Andriy, Hinton Geoffrey E. A Scalable Hierarchical Distributed Language Model // NeurIPS. — 2008.*



Hierarchical softmax


- We still have a **fully connected neural network**
- The only thing that differs is **softmax** calculation
- To calculate loss we don't need the whole vector of probabilities
- We only need the values in positions of **words in the context window**:

$$\sum_{i=1}^N \sum_{j=-C, j \neq 0}^C \log p(w_{i+j}|w_i) \rightarrow \max_{W,V}$$



Hierarchical softmax


- Let's consider **skip-gram** model
- Hidden state after first layer $h = W^T x$
- Let's change the second layer
into a binary tree (for example, Huffman tree)
- Assign every leaf one word from vocabulary
- Assign every internal node a vector of **k** weights



Hierarchical softmax

- Suppose we want to obtain a probability of a word $w = \text{"tea"}$
- Probability of paths (left and right) in current node:

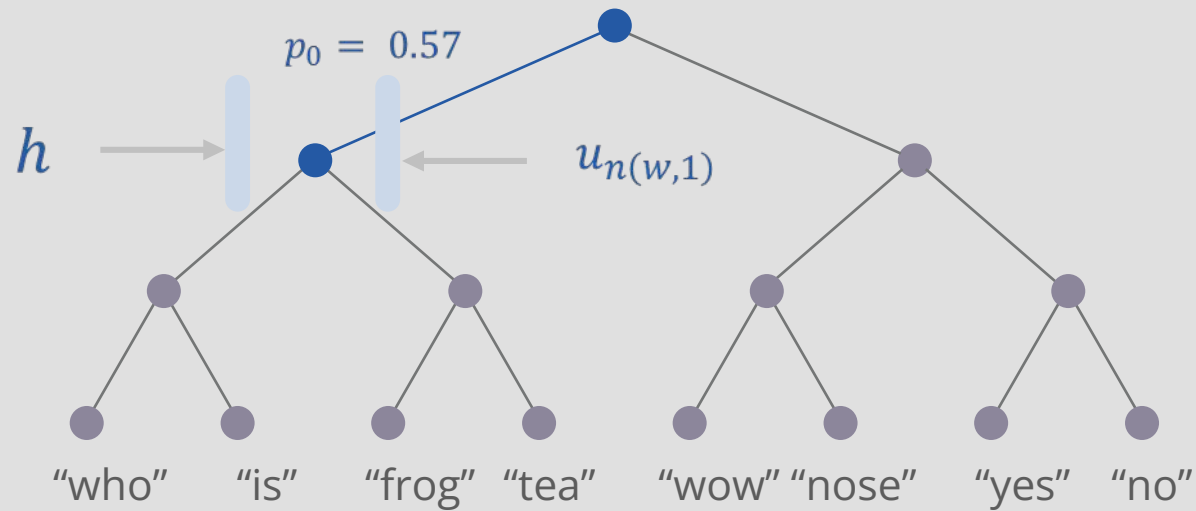
$$p_0 = \sigma(u_{n(w,0)}^T h) \quad p_1 = \sigma(-u_{n(w,0)}^T h)$$



Hierarchical softmax


- Do the same in the next node:

$$p_1 = \sigma(-u_{n(w,1)}^T h)$$



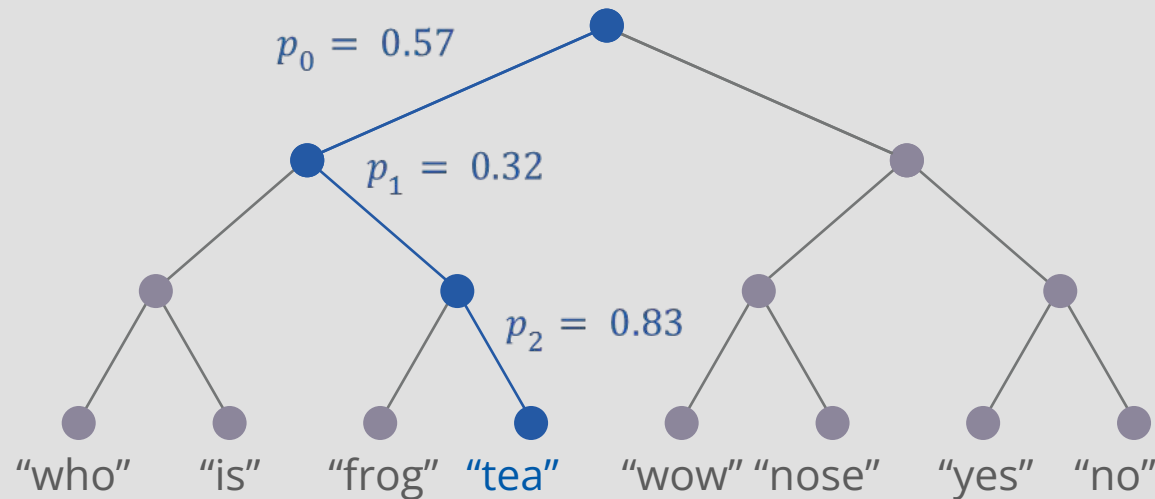
Hierarchical softmax

$$p_2 = \sigma(-u_{n(w,2)}^T h)$$



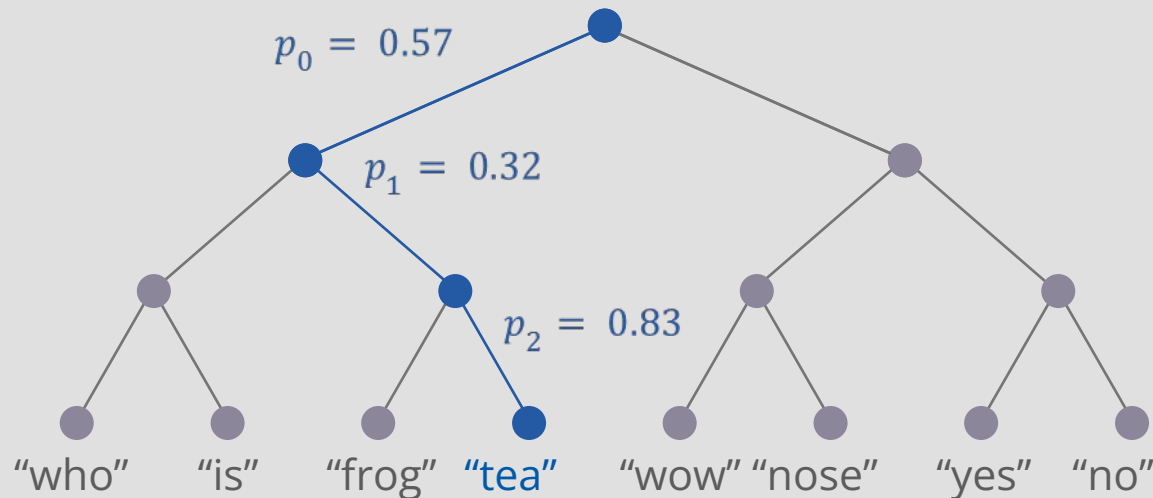
Hierarchical softmax

- In the end, we are at leaf with the word tea
- Every step i had a probability p_i



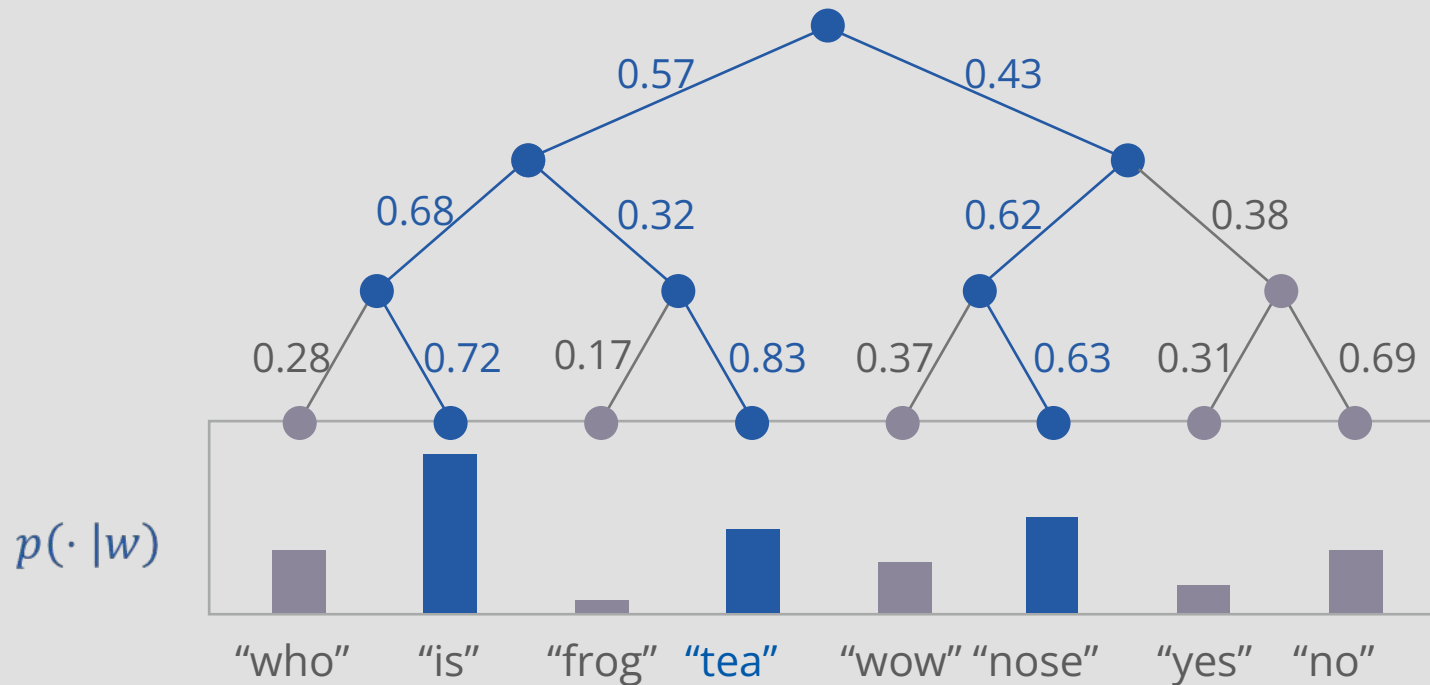
Hierarchical softmax

- In the end, we are at leaf with the word tea
- Every step i had a probability p_i
- The final probability of path is $\prod_i p_i$



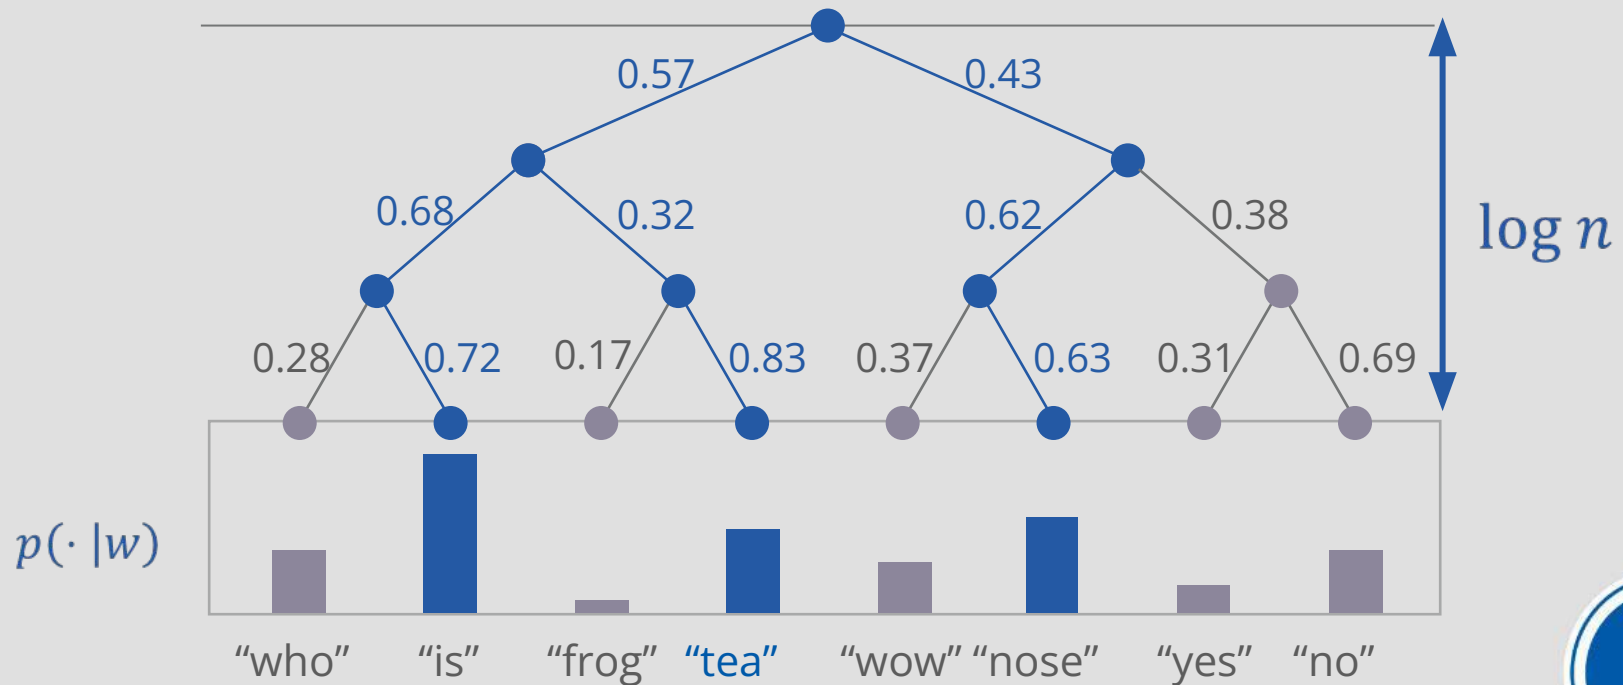
Hierarchical softmax

- If we do the procedure 2C times we obtain all probabilities required to calculate loss



Hierarchical softmax

- If we do the procedure $2C$ times we obtain all probabilities required to calculate loss
- The complexity of current calculation - $O(\log n)$



All in all

Before:

- Obtain h from the first layer
- Multiply h with second layer matrix V
- Apply softmax
- Use only $2C$ probabilities from the softmax result



All in all

Before:

- Obtain h from the first layer
- Multiply h with second layer matrix V
- Apply softmax
- Use only $2C$ probabilities from the softmax result

Now:

- Obtain h from the first layer
- $2C$ times go through the tree
- Obtain only essential probabilities



Skip-gram negative sampling

- For skip-gram model there is another popular training optimization method



Skip-gram negative sampling

- For **skip-gram** model there is another popular training optimization method
- We change the formulation of the problem and loss function
- We will solve **binary classification problem**



Skip-gram negative sampling

- For skip-gram model there is another popular training optimization method
- We change the formulation of the problem and loss function
- We will solve binary classification problem
- Object - pair of words (w, s)
- Class 1: word s belongs to context w
- Class 2: word s doesn't belong to context w
- For each word w we compare trainable vector v_w which will be the sought one



What are the advantages?

- Training the model for each input object requires an update of all weights of the input layer
- Softmax in classic approach leads to an update of all weights in all layers
- In new approach we only update the weights of the layers, that are involved in the current iteration of training



Skip-gram negative sampling

- Class probability is simulated with sigmoid:

$$p(1|(w, s)) = \frac{1}{1 + \exp(-v_w^T v_s)} = \sigma(v_w^T v_s)$$



Skip-gram negative sampling

- Class probability is simulated with sigmoid:

$$p(1|(w, s)) = \frac{1}{1 + \exp(-v_w^T v_s)} = \sigma(v_w^T v_s)$$

- Let D_1 be a subset of pairs (w, s) where s belongs to context w
- Let D_2 be a subset of all other possible pairs
- Then the likelihood function is:

$$L = \sum_{(w,s) \in D_1} \log(\sigma(v_w^T v_s)) + \sum_{(w,s) \in D_2} \log(\sigma(-v_w^T v_s))$$



Skip-gram negative sampling

If we optimize this function, we solve the task!



Skip-gram negative sampling

If we optimize this function, we solve the task!

But:

- We need examples of pairs
- D_1 can be obtained from data, while D_2 is not presented in data as is



Skip-gram negative sampling

If we optimize this function, we solve the task!

But:

- We need examples of pairs
- D_1 can be obtained from data, while D_2 is not presented in data as is

Solution: generate negative samples by sampling random pairs of words on each training step



Main conclusions

- Canonical word2vec implementation scales poorly on dictionary and corpus volume
- Main difficulty is the second layer and softmax calculation
- There are several training optimization methods, the main ones are hierarchical softmax and negative sampling




GloVe model



GloVe (Global Vectors)

- Construct a matrix $\mathbf{X} \in \mathbb{R}^{n \times n}$
- Every column and row correspond to a word from a dictionary




- Pennington Jeffrey, Socher Richard, Manning Christopher D. Glove: Global Vectors for Word Representation. // EMNLP. — Vol. 14. — 2014. — Pp. 1532–1543.




GloVe (Global Vectors)

- Construct a matrix $\mathbf{X} \in \mathbb{R}^{n \times n}$
- Every column and row correspond to a word from a dictionary
- x_{ij} - number of times word i occurs in context of word j




GloVe (Global Vectors)

- Construct a matrix $\mathbf{X} \in \mathbb{R}^{n \times n}$
- Every column and row correspond to a word from a dictionary
- x_{ij} - number of times word i occurs in context of word j
- $X_i = \sum_j x_{ij}$ - sum of elements of row i



GloVe (Global Vectors)

- Construct a matrix $\mathbf{X} \in \mathbb{R}^{n \times n}$
- Every column and row correspond to a word from a dictionary
- x_{ij} - number of times word i occurs in context of word j
- $X_i = \sum_j x_{ij}$ - sum of elements of row i
- $P_{ij} = \frac{x_{ij}}{X_i}$ - probability of word j to occur in context of word i



GloVe (Global Vectors)

- Assume we know vector representation v_i for every word i
- Also assume that we know all x_{ij}



GloVe (Global Vectors)

- Assume we know vector representation v_i for every word i
- Also assume that we know all x_{ij}
- Define function $F\left(f(v_i, v_j, v_k)\right) = \frac{P_{ik}}{P_{jk}}$
- This function F shows which one of words i and j is more likely to occur in context of word k
- Function f is some function from input to real number



GloVe (Global Vectors)

- We need F to satisfy

$$F((v_i - v_j)^T v_k) = \frac{F(v_i^T v_k)}{F(v_j^T v_k)} = \frac{P_{ik}}{P_{ij}}$$

- F can be $\exp(x)$



GloVe (Global Vectors)

- So

$$F(x) = \exp(x)$$

$$P_{ij} = \frac{x_{ij}}{X_i}$$



GloVe (Global Vectors)

- So

$$F(x) = \exp(x)$$

$$P_{ij} = \frac{x_{ij}}{X_i}$$

- Rewrite

$$F(v_i^T v_k) = P_{ik}$$

$$v_i^T v_k = \log(x_{ik}) - \log(X_i)$$



GloVe (Global Vectors)

- So

$$F(x) = \exp(x)$$

$$P_{ij} = \frac{x_{ij}}{X_i}$$

- Rewrite

$$F(v_i^T v_k) = P_{ik}$$

$$v_i^T v_k = \log(x_{ik}) - \log(X_i)$$

- Now we remember that we only have x_{ij}



GloVe (Global Vectors)

- we want

$$v_i^T v_k = \log(x_{ik}) - \log(X_i)$$

- Therefore we rewrite

$$\sum_i \sum_k F(x_{ik}) (v_i^T v_k + b_i + b_k - \log(x_{ik}))^2 \rightarrow \min_{v_i, b_i, i \in \{1, n\}},$$

$$\log(X_i) = b_i + b_k$$

and obtain word embeddings.



Main conclusions

- Another example on an approach based on frequencies
- In practice, it works quite similar to word2vec
- There are many pre-trained GloVe models



FastText model, Hashing Trick



Word2vec and GloVe problems:

- Problem 1: Out-of-Vocabulary (OOV) words




Word2vec and GloVe problems:

- **Problem 1:** Out-of-Vocabulary (OOV) words
We train our model on corpus V
Then we try to obtain a vector for a new word w
But there is no embedding for it



Word2vec and GloVe problems:

- **Problem 1:** Out-of-Vocabulary (OOV) words
We train our model on corpus V
Then we try to obtain a vector for a new word w
But there is no embedding for it



Word2vec and GloVe problems:

- Problem 2: Lack of consideration of morphology
- Suppose we train a model for a language with rich morphology (for example, russian):



Word2vec and GloVe problems:

- **Problem 2:** Lack of consideration of morphology
- Suppose we train a model for a language with rich morphology (for example, russian):
- We get a new embedding for each word
- **As a result:**
 - Many similar embeddings (and redundant memory)
 - Less samples for each word in training data

яблоко
яблока
яблоку
...
яблоками
яблоках




Character embeddings

- We can try to get embeddings for a smaller piece of language
- For example, **for each character**
- All training methods remain the same
- Word embedding can be obtained by averaging character embedding



Character embeddings

- We can try to get embeddings for a smaller piece of language
- For example, **for each character**
- All training methods remain the same
- Word embedding can be obtained by averaging character embedding



Character n-grams embeddings

- Instead of characters we can use **character n-grams**
- It improves the quality of the resulting word embeddings significantly



Character n-grams embeddings

- Instead of characters we can use **character n-grams**
- It improves the quality of the resulting word embeddings significantly

N-grams for a word «doctor»

N=3: ^do, doc, oct, cto, tor, or\$

N=4: ^doc, doct, octo, octor, tor\$

N=5: ^doct, docto, octor, ctor\$



Character n-grams embeddings

- Instead of characters we can use **character n-grams**
- It improves the quality of the resulting word embeddings significantly

N-grams for a word «doctor»

N=3: ^do, doc, oct, cto, tor, or\$

N=4: ^doc, doct, octo, octor, tor\$

N=5: ^doct, docto, octor, ctor\$

- Word embeddings are obtained by averaging



Word2vec and GloVe problems

Improvements:

- OOV problem solved
- Morphology problems also solved

But:

- There can be even more sequences of characters than there are different variants of words
- Tens of millions of vectors may simply not fit into RAM



Hash-functions and hash-tables

- String hash-function converts a string into a number



Hash-functions and hash-tables

- String hash-function converts a string into a number
- Function requirements:
 - The range of values is limited by an interval
 - Function values are distributed approximately uniformly over the interval



Hash-functions and hash-tables

- **String hash-function** converts a string into a number
- Function requirements:
 - The range of values is limited by an interval
 - Function values are distributed approximately uniformly over the interval
- **Hash-table** — an array of values + hash-function that transforms input string into array's indices



Hashing Trick

- We fix the maximum number of vectors that we want to train



Hashing Trick

- We fix the maximum number of vectors that we want to train
- Match them a hash-table



Hashing Trick

- We fix the maximum number of vectors that we want to train
- Match them a hash-table
- All n-gramm char-embeddings are distributed over the values of array



Hashing Trick

- We fix the maximum number of vectors that we want to train
- Match them a hash-table
- All n-gramm char-embeddings are distributed over the values of array
- Several n-grams use the same embedding



FastText

- Package that trains word embeddings
 - Uses **CBOW / skip-gram** both for words and character n-grams
 - Optimizes RAM consumption by **hashing trick**
 - Parallels training process on **CPU**
-
- *Enriching Word Vectors with Subword Information / Piotr Bojanowski, Edouard Grave, Armand Joulin, Tomas Mikolov // Transactions of the Association for Computational Linguistics. — 2017. — Vol. 5. — Pp. 135–146.*



Main conclusions

- Classic word2vec models work poorly with OOV words and morphology
- Models that operate with word fragments can solve these problems
- FastText package allows to train such models efficiently on CPU



Word embeddings evaluation



Types of metrics

- Word embeddings' quality can be measured by internal and external criteria



Types of metrics

- Word embeddings' quality can be measured by internal and external criteria
- Internal:
 - Quality of similar words search
 - Quality of analogies solving



Types of metrics

- Word embeddings' quality can be measured by internal and external criteria
- **Internal:**
 - Quality of similar words search
 - Quality of analogies solving
- **External:**
 - Quality of the final problem solution (the problem you use word embedding in)



Similar words search

- You have a word embedding model
- You have a corpus with human-evaluated semantic similarity between words



Similar words search

- You have a word embedding model
- You have a corpus with human-evaluated semantic similarity between words
- Calculate cosine similarity between word embeddings:

REMINDER

$$\cos(\theta) = \frac{p \cdot q}{\|p\| \|q\|} = \frac{\sqrt{\sum_{i=1}^n p_i q_i}}{\sqrt{\sum_{i=1}^n p_i^2} \sqrt{\sum_{i=1}^n q_i^2}}$$



Similar words search

- You have a word embedding model
- You have a corpus with human-evaluated semantic similarity between words
- Calculate cosine similarity between word embeddings:

REMINDER

$$\cos(\theta) = \frac{p \cdot q}{\|p\| \|q\|} = \frac{\sqrt{\sum_{i=1}^n p_i q_i}}{\sqrt{\sum_{i=1}^n p_i^2} \sqrt{\sum_{i=1}^n q_i^2}}$$

- Check that it correlates with human evaluation

$\cos(\text{абрикос}, \text{персик}) > \cos(\text{абрикос}, \text{маска})$



Analogies solving

- You have triplets of words a, a^*, b
- Words a, a^* have some kind of relation

- $a = \text{читать}, a^* = \text{чтение}$
 $b = \text{петь}$



Analogies solving

- You have triplets of words a, a^*, b
- Words a, a^* have some kind of relation
- We want to find a word b^* that has the same kind of relation with the word b

- $a = \text{читать}, a^* = \text{чтение}$
 $b = \text{петь}, b^* = \text{пение}$



Analogies solving


- You have triplets of words a, a^*, b
- Words a, a^* have some kind of relation
- We want to find a word b^* that has the same kind of relation with the word b

$$a^* - a + b$$



- $a = \text{читать}, a^* = \text{чтение}$
 $b = \text{петь}, b^* = \text{пение}$
 $a^* - a + b \approx b^*$
 $\text{чтение} - \text{читать} + \text{петь} \approx \text{пение}$





Analogies solving



Analogies solving



Analogies solving



External metrics for word embeddings models

- We are not interested in evaluation of word embedding themselves



External metrics for word embeddings models

- We are not interested in evaluation of word embedding themselves
- We solve some task with word embedding models and track the metrics of the solution



External metrics for word embeddings models

- We are not interested in evaluation of word embedding themselves
- We solve some task with word embedding models and track the metrics of the solution
- For example, we can solve some classification problem with TF-IDF vectors of FastText



External metrics for word embeddings models

- We are not interested in evaluation of word embedding themselves
- We solve some task with word embedding models and track the metrics of the solution
- For example, we can solve some classification problem with TF-IDF vectors of FastText
- Mutual quality may differ for different tasks and datasets



Main conclusions

- The quality of word embeddings can be measured by different methods
- Internal criteria measure the models' quality in terms of their internal properties
- External criteria are more abstract and focus on the final problem where the word embedding model is applied

